

Introduction to Apache Kafka

Martin Brucker

June 20. 2024

Table of Contents

- 1 Introduction
- 2 Distributed Message Brokers and Log Systems
- 3 IT Architecture of Kafka
- 4 Client Libraries
- 5 Typical Use Cases
- 6 Installation and Configuration
- 7 Programming with Kafka
- 8 Conclusion

Introduction

- **Apache Kafka**: A distributed streaming platform designed for high-throughput, fault-tolerant, and real-time data processing.
- Combines the features of both a **messaging system** and a **log system**.
- Originally developed by LinkedIn and later open-sourced.

Distributed Message Brokers and Log Systems

- **Message Brokers:**
 - Queues and topics for asynchronous communication.
 - Point-to-point and publish-subscribe models.
- **Log Systems:**
 - **Distributed commit log:** Stores all published messages.
 - **Append-only storage:** Immutable logs for durability.

IT Architecture of Kafka

- **Topics:** Logical channels for data streams.
- **Partitions:** Segments of topics stored on different servers.
- **Producers:** Data publishers that write to topics.
- **Consumers:** Data subscribers that read from topics.

- **Java:** Official client library (highly recommended).
- **Python:** Confluent's Python client (for Python enthusiasts).
- **Other languages:** Community-supported libraries (e.g., Go, .NET).

Typical Use Cases

- **Real-time analytics:** Process and analyze data as it arrives.
- **Log aggregation:** Centralize logs from various services.
- **Event sourcing:** Capture and replay events for stateful systems.
- **Stream processing:** Transform and enrich data streams.

Installation and Configuration

- **Download and install Kafka:** Obtain the Kafka distribution.
- **Configure Kafka properties:**
 - Specify server settings (e.g., broker ID, port).
 - Define topics and their partitions.
- **Start Kafka server and Zookeeper ensemble:**
 - Kafka relies on Zookeeper for coordination.
 - Start both services to enable Kafka functionality.

Programming with Kafka

- **Produce messages** using Kafka producers.
- **Consume messages** using Kafka consumers.
- Handle message **serialization and deserialization**.

Conclusion

- Apache Kafka is a powerful tool for building scalable and reliable data pipelines.
- Dive deeper into its documentation and explore advanced usage scenarios.